

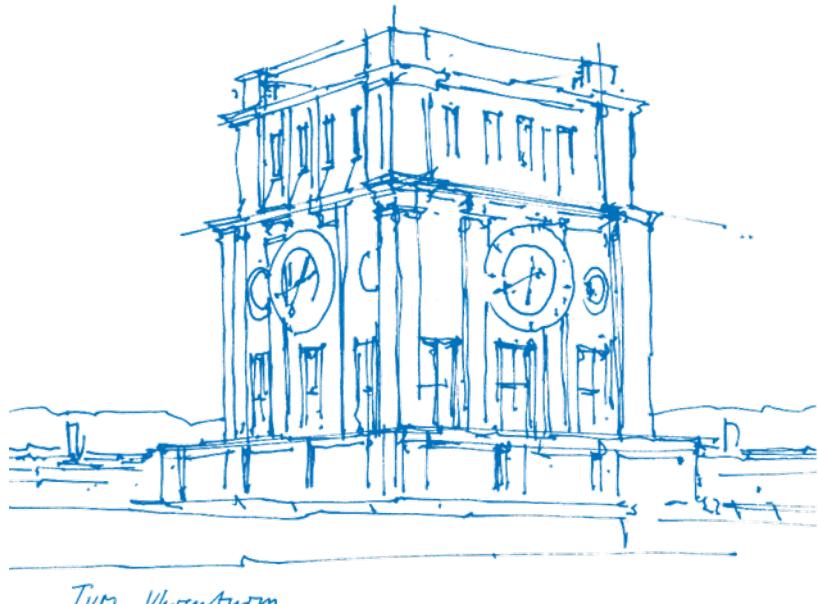
Tutorial 01¹ - Sequential Programming

Bengisu Elis, M.Sc.

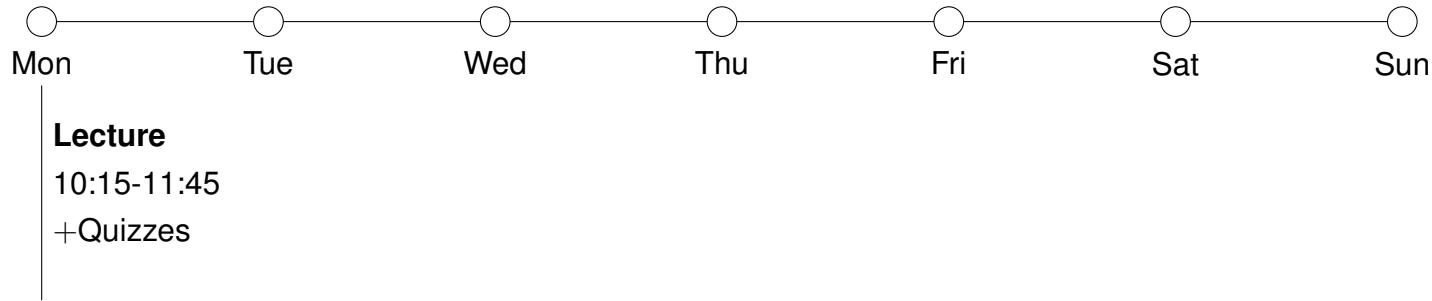
Chair for Computer Architecture and Parallel Systems (CAPS)

Technical University Munich

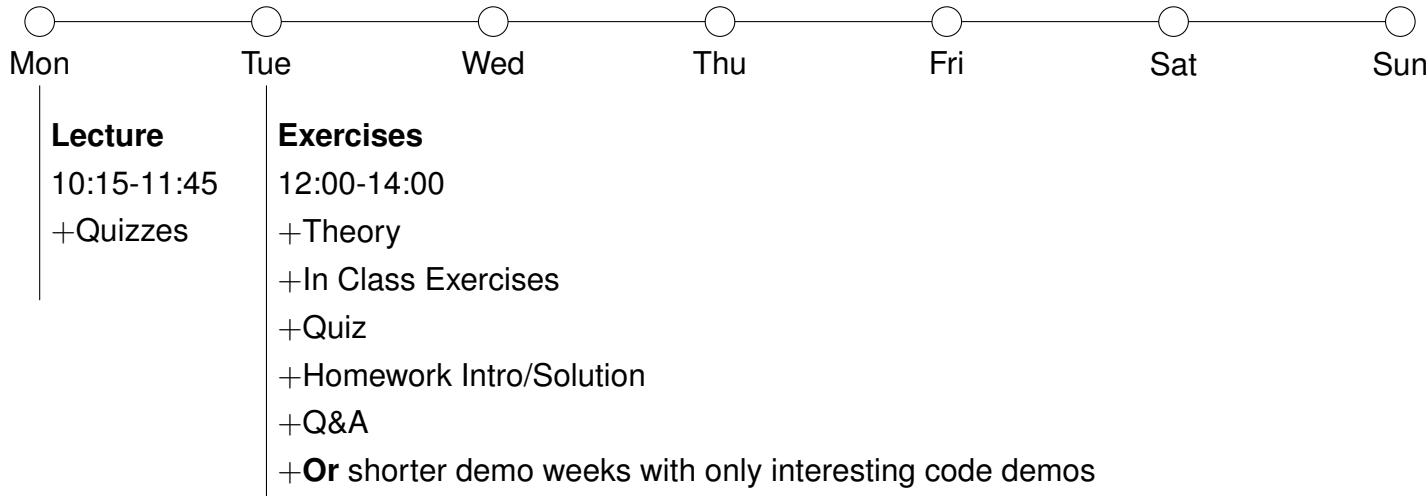
May 2, 2022



¹Courtesy of Vincent Bode

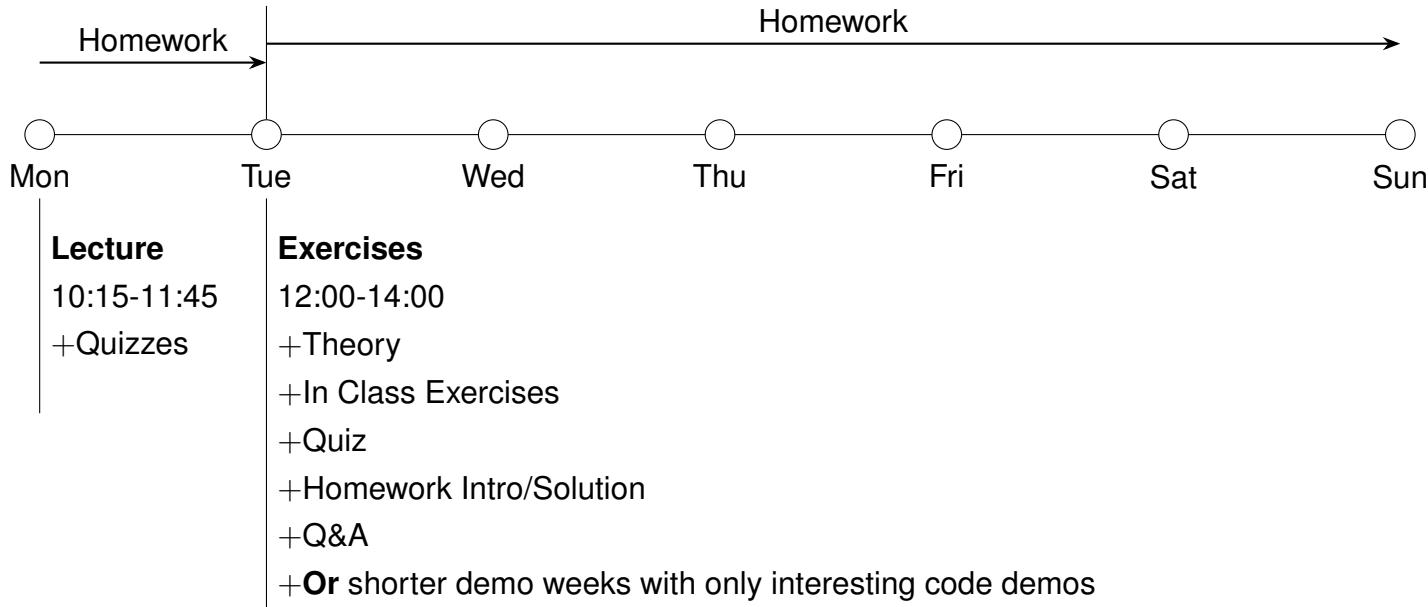


Organization



Organization

Parallel Programming 2022
Tutorial 01





Lecture + Exercise Content

www.moodle.tum.de/course/view.php?id=63406



Programming Assignments

parprog.caps.in.tum.de

<https://gitlab.lrz.de/lrr-tum/teaching/parprog/ss2022/published-assignments.git>



Questions and Chat

chat.tum.de/channel/ParProg22

The Team



Prof. Martin Schulz
Lecturer



Dennis-Florian Herr
Teaching Assistant



Bengisu Elis
Teaching Assistant

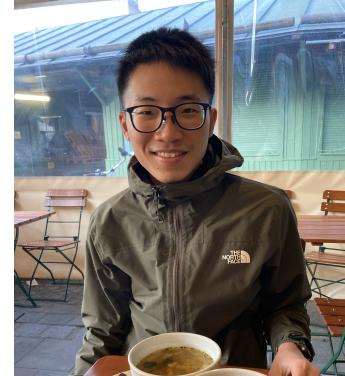
The Team: Tutors



Nathan Mateo Marin Jimenez



Yakub Koray Budanaz



Hua-Ming Huang



Fauzan Ramadhan

Assignments

In class exercises

- Entirely optional
- Relevant to homework
- Can also be solved after session

Assignments

In class exercises

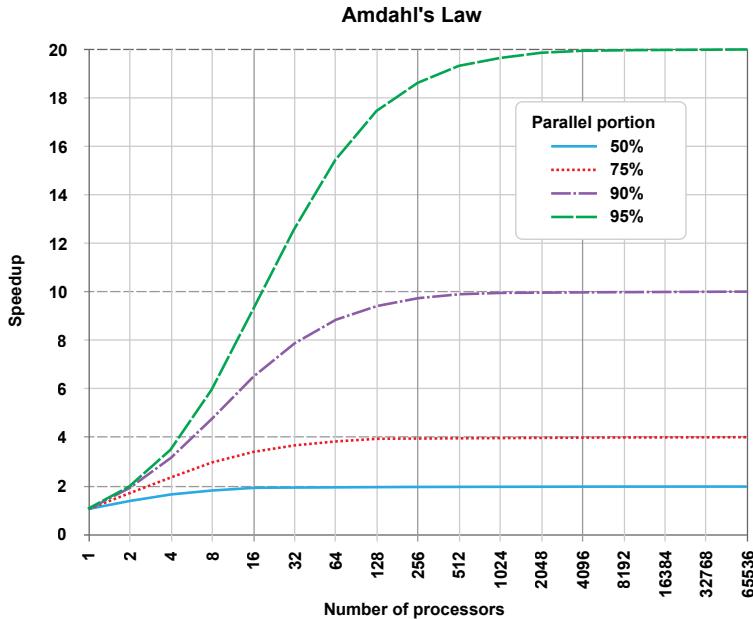
- Entirely optional
- Relevant to homework
- Can also be solved after session

Homework

- In teams with up to 3 members
- Each assignment is doable by 1 person
- Approx. 1/week with 1 week of time to solve
- Grade Bonus
 - 8 homeworks in total. (maybe subject to change)
 - Bonus earned if 75% of all homeworks successful (6 out of 8)
 - Bonus = grade enhancement of exam result by 0.3.
 - Only applies for a passed exam (original grade ≤ 4.0)
- Online submission with almost instant feedback
- We check: correctness, speedup, and plagiarism
- Solutions discussed in next exercise session

Theory

Amdahl's Law & Parallel Efficiency



$$\text{Speedup} = \frac{1}{(1-p) + p/N} \quad (1)$$

p = parallelizable part of the program
 N = number of processors

Figure 1: Well life is not perfect ! ([Source](#)). But rarely it can be optimized to be perfect.

Definition

There is a data dependence from a source statement to a sink statement (sink depends on source), if and only if

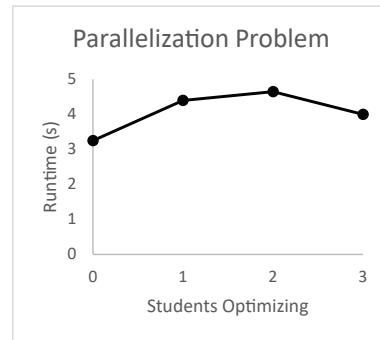
1. both statements access the same memory location and at least one of them writes to it.
2. there is a feasible run-time execution path from the source to the sink.

Moodle Quiz

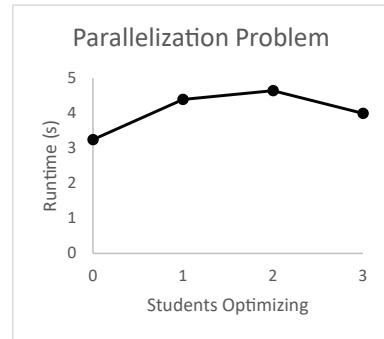
<https://www.moodle.tum.de/mod/quiz/view.php?id=2143814>

15 minutes reserved for Quiz. Post questions to rocket chat or ask in-class.

Original Runtime	Students Optimizing	Runtime Reduction	Optimized Runtime	Students Parallelizing	Parallel Code	Parallel Runtime	Sequential Code	Sequential Runtime	Total Runtime
10s	0	0%	10s	3	90%	2.25s	10%	1.00s	3.25s
10s	1	20%	8s	2	60%	1.20s	40%	3.20s	4.40s
10s	2	40%	6s	1	30%	0.45s	70%	4.20s	4.65s
10s	3	60%	4s	0	0%	0.00s	100%	4.00s	4.00s



Original Runtime	Students Optimizing	Runtime Reduction	Optimized Runtime	Students Parallelizing	Parallel Code	Parallel Runtime	Sequential Code	Sequential Runtime	Total Runtime
10s	0	0%	10s	3	90%	2.25s	10%	1.00s	3.25s
10s	1	20%	8s	2	60%	1.20s	40%	3.20s	4.40s
10s	2	40%	6s	1	30%	0.45s	70%	4.20s	4.65s
10s	3	60%	4s	0	0%	0.00s	100%	4.00s	4.00s



t = original runtime ; z = students parallelizing ; x = students optimizing
new runtime:

$$[t(1 - 0.3z) + \frac{t \cdot 0.3z}{4}] \cdot (1 - 0.2x) \quad (2)$$

$$x + z = 3 \text{ for } z \geq 0 \text{ and } x \geq 0$$

In Class Exercise 1

Where to find the exercise ?

- Go to the following repository to get the exercise:

<https://gitlab.lrz.de/lrr-tum/teaching/parprog/ss2022/published-assignments>

- Use git to clone the exercise to your local machine:

cd your_folder

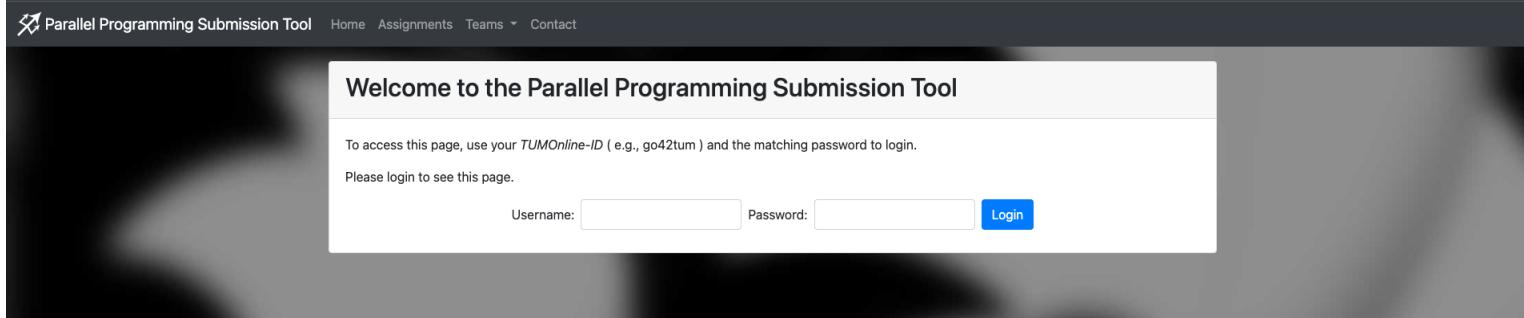
git clone <https://gitlab.lrz.de/lrr-tum/teaching/parprog/ss2022/published-assignments.git>

- You can pull from this repository every time a new exercise is issued

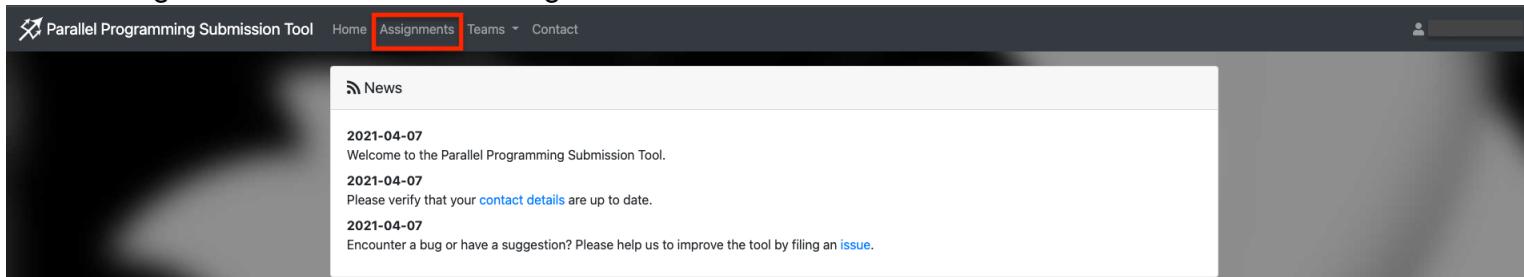
- Go to folder Week_1 for this week's task, you can also find a README.md there

How to submit the solution ?

- Go to : <https://parprog.caps.in.tum.de/Submission/login?next=/Submission/home>
- Log-in with TUM Online ID



- Click "Assignments" section in the navigation bar:



- Upload your code either by copying the code into the textbox or uploading the "student_submission.cpp" file
- After submission you will see your speed up on the same page.

Video Resource

In Class Exercise 1

- **Task 1: Optimize the decryption process**

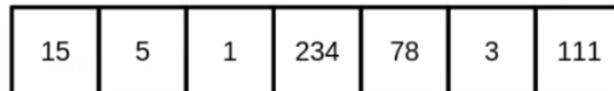
- decrypt a message `DecryptedMessage` using dictionary where the keys are stored in `keys` and the values are stored in `values` .

DecryptedMessage

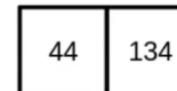


...

keys

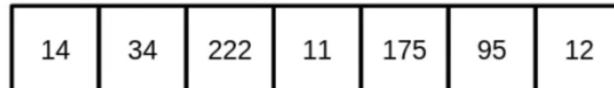


...



225

values

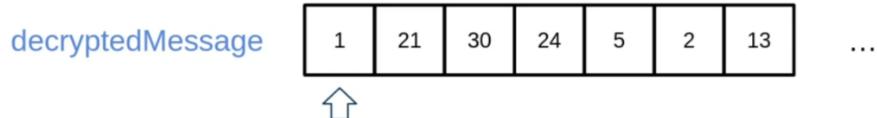


...

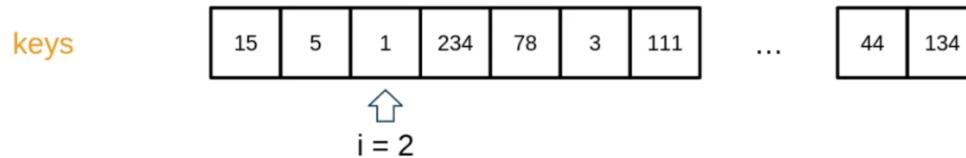


225

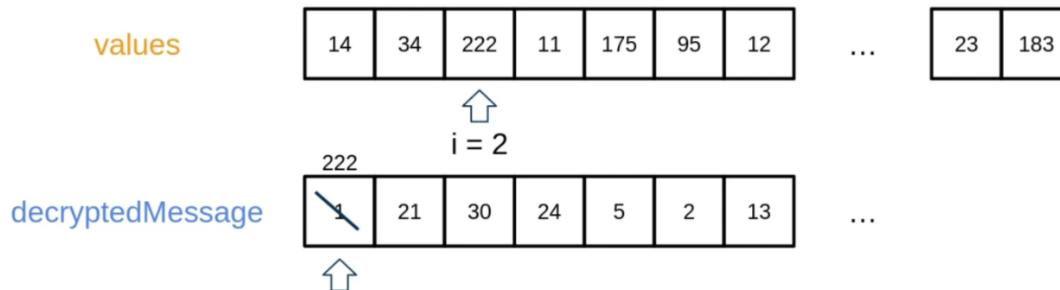
In Class Exercise 1



1. For each character in the message, search the character in **keys**, get the index i in **keys**



2. Use index i to find a character in **values** and replace the one in the message



In Class Exercise 1

decryptedMessage

1	21	30	24	5	2	13
...						

1. For each character in the message, search the character in **keys**, get the index i in **keys**

keys

15	5	1	234	78	3	111
...						

$i = 2$

2. Use index i to find a character in **values** and replace the one in the message

values

14	34	222	11	175	95	12
...						

$i = 2$

decryptedMessage

222	21	30	24	5	2	13
...						

\uparrow

$\times 50.000$

In Class Exercise 1



10 mins reserved for In Class Exercise 1

Listing 1: Initial Code

```

1 int index = -1;
2 for (unsigned int i = 0; i < STRING_LEN; ++i){
3     for (int k = 0; k < UNIQUE_CHARACTERS; ++k){
4         if (decryptedMessage[i] == keys[k]){
5             index = k;
6         }
7     }
8     decryptedMessage[i] = values[index];
9 }
```

Listing 2: Optimized Code

```

1 int dict[UNIQUE_CHARACTERS];
2 for (unsigned int i = 0; i < UNIQUE_CHARACTERS; ++i){
3     uint8_t key = keys[i];
4     dict[key] = values[i];
5 }
6
7 // use the map instead of searching in originalCharacter
8 for (unsigned int i = 0; i < STRING_LEN; ++i){
9     decryptedMessage[i] = dict[decryptedMessage[i]];
10 }
```

Homework Assignment 1

Go find a team!

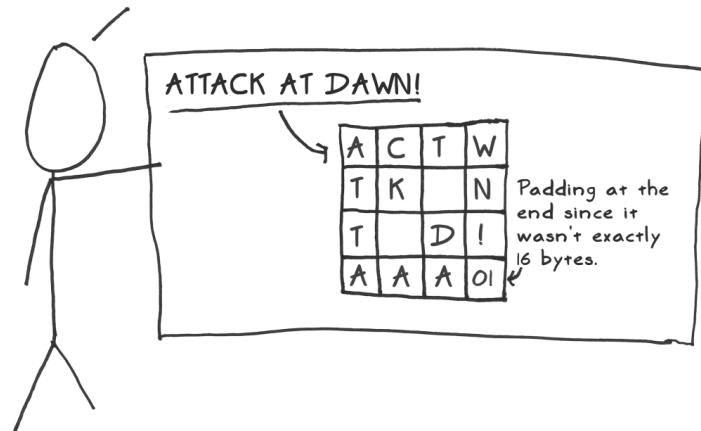
Up to 3 students per team.

<https://www.moodle.tum.de/mod/forum/view.php?id=2149207>

Explanation: VV-AES

Step 1:

*I take your data and load it
into this 4x4 square.**



** This is the 'state matrix' that I carry with me at all times.*

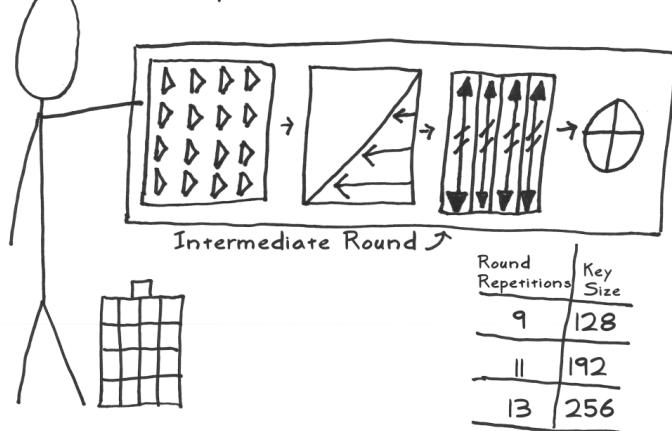
2

²Courtesy of Jeff Moser's [A Stick Figure Guide to the Advanced Encryption Standard \(AES\)](#)

Explanation: VV-AES

Step 2:

Next, I start the intermediate rounds. A round is just a series of steps I repeat several times. The number of repetitions depends on the size of the key.



3

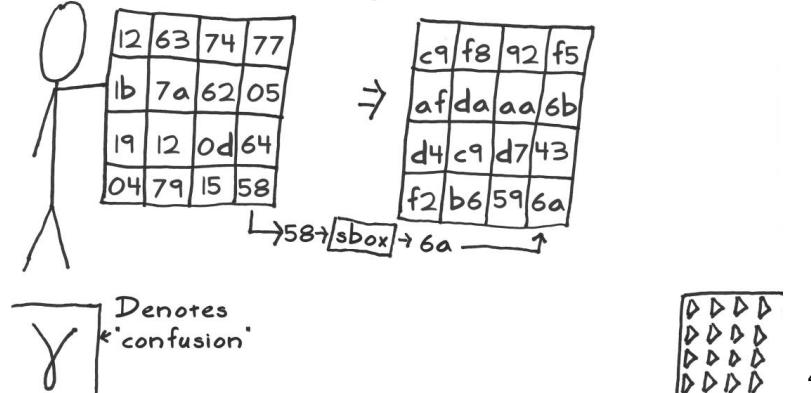
³Courtesy of Jeff Moser's [A Stick Figure Guide to the Advanced Encryption Standard \(AES\)](#)

Explanation: VV-AES

Step 2.1:

Applying Confusion: Substitute Bytes

I use confusion (Big Idea #1) to obscure the relationship of each byte. I put each byte into a substitution box (sbox), which will map it to a different byte:



⁴Courtesy of Jeff Moser's [A Stick Figure Guide to the Advanced Encryption Standard \(AES\)](#)

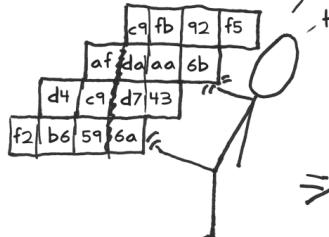
Explanation: VV-AES

Step 2.2:

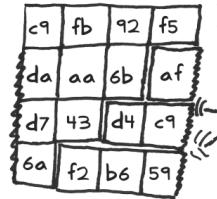
Applying Diffusion, Part 1: Shift Rows

Next I shift the rows to the left

Hiiiiii yaah!



...and then wrap them around the other side

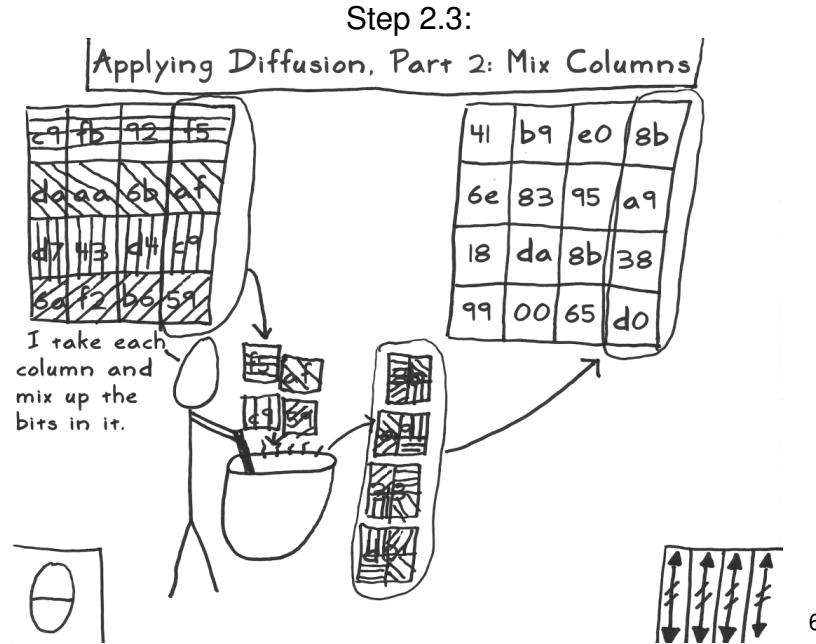


5

Denotes
permutation

⁵Courtesy of Jeff Moser's [A Stick Figure Guide to the Advanced Encryption Standard \(AES\)](#)

Explanation: VV-AES



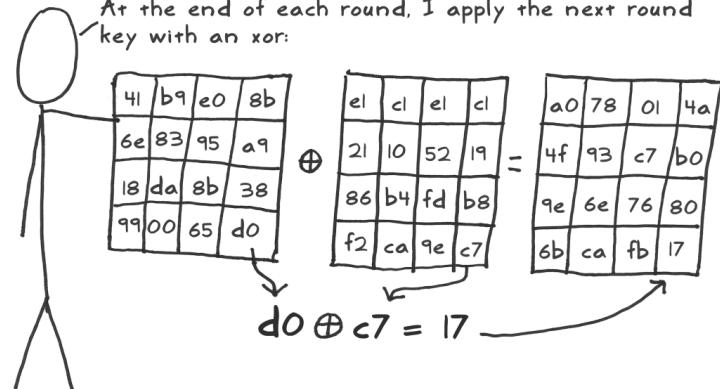
⁶Courtesy of Jeff Moser's [A Stick Figure Guide to the Advanced Encryption Standard \(AES\)](#)

Explanation: VV-AES

Step 2.4:

Applying Key Secrecy: Add Round Key

At the end of each round, I apply the next round key with an xor:



7

⁷Courtesy of Jeff Moser's [A Stick Figure Guide to the Advanced Encryption Standard \(AES\)](#)

Constraints

- Only a single Thread.
- Locked in a container.
- Eventually killed when it takes too long to do its job.
- Reads input from a file called *stdin*.
- Writes output to a file called *stdout*.
- Can file complaints as they like to *stderr*.
- Can only remember $\sim 10^9$ pieces of information at a time without writing them down.
- Lacks access to the network.
- Any paper trail is burned after every batch of encrypting.

Your job: Optimize the Single threaded Workflow

- Speedup of 8x in relation to the provided implementation.
- Note: VV-AES is a slightly simplified version of AES.
 - More information on real AES <http://www.moserware.com/2009/09/stick-figure-guide-to-advanced.html>.

Typical performance problems:

- Frequent memory allocation/deallocation
- Unnecessary copying or conversion of data
- Lack of cache awareness
- **Bad algorithms**
- **Reinventing the wheel**

Developing

- Copy “sequential_implementation.cpp” to “student_submission.cpp”
- Write your code in “student_submission.cpp”

Build the program

- Makefile:
\$> make

Usage of the program

- Sequential:
\$> echo \$RANDOM | ./sequential_implementation
- Your solution:
\$> echo \$RANDOM | ./student_submission
- These programs expect a number to seed random problem generation on `stdin`. Replace the random number with a specific number when you want to re-run a specific problem.

No Linux? No Problem.

You have several options for getting access to a linux environment.

- Install linux in a virtual machine (e.g. VirtualBox)
 - Don't forget to assign multiple cores to the virtual machine
- Use WSL (Windows Subsystem for Linux) on Windows 10
- Use Rechnerhalle
 - Accessible at the workstations or remotely.
 - Remote ssh access: ssh <rbg-id>@lxhalle.in.tum.de
 - Your RBG-Id is the part in front of your @in.tum.de or @ma.tum.de email address.
- Ask the tutors; they will be more than happy to help you.

Covered today:

- Amdahl's Law
- Speedup & Parallel Efficiency
- Race Conditions

Questions

This slide is intentionally left blank.